

Formal Lab Report 1: Numerical Errors and Simple Numerical Integration

Due Sep 30 2024

Background

Solving integrals numerically: Lecture notes, as well as Sections 5.1 – 5.3 of the text, introduce the Trapezoidal Rule and Simpson’s Rule. The online resource for the textbook provides the Python program `trapezoidal.py`, which you are free to copy-and-paste directly or to adapt for this lab.

Standard deviation calculations: To calculate the sample mean \bar{x} and standard deviation σ of a sequence $\{x_1, \dots, x_n\}$ using the standard formulas

$$\bar{x} \equiv \frac{1}{n} \sum_{i=1}^n x_i \quad \text{and} \quad \sigma \equiv \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2} \quad (1)$$

requires two passes through the data: the first to calculate the mean and the second to compute the standard deviation (by subtracting the mean in the term $(x_i - \bar{x})$ before squaring it). There is an alternative, mathematically equivalent, formula for the standard deviation:

$$\sigma \equiv \sqrt{\frac{1}{n-1} \left(\sum_{i=1}^n x_i^2 - n\bar{x}^2 \right)}, \quad (2)$$

This might seem preferable because, if you think about it, you can calculate σ in eqn. (2) with only a single pass through the data. This means, for example, that you could calculate these statistics on a live incoming data stream as it updates. However, the one-pass method has numerical issues that you will investigate in Q1.

The function `numpy.std(array, ddof=1)` also implements eqn. (1). It has been highly optimized. You can think of as a “correct” calculation.

Numerical (roundoff) error: Read Section 4.2 of Newman’s textbook, which discusses characteristics of machine error. One important point is that you can treat errors on numerical calculations as random and independent. (This is a little confusing because you will find

that errors on a given computer are often *reproducible*: they'll come out the same if you do the calculation the same way multiple times on your computer. But there is typically no way to predict what this error value is, i.e. it could be different on different computers, or even if you do a software update on your computer.) As a result of this, you can use standard error statistics, as in experimental physics, to figure out how error propagates through numerical calculations. This results in expressions like (4.7) in the text, describing the error on the sum (series) of N terms $x = x_1 + \dots + x_N$:

$$\sigma = C\sqrt{N}\sqrt{\overline{x^2}}, \quad (3)$$

where $C \approx 10^{-16}$ is the *error constant* defined on p.130 and the overbar indicates a mean value. This means that the more numbers we include in a given series, the larger the error (by $O(\sqrt{N})$). Even if the mean error is small compared to the individual terms the *fractional error*

$$\frac{\sigma}{\sum_i x_i} = \frac{C}{\sqrt{N}} \frac{\sqrt{\overline{x^2}}}{\bar{x}} \quad (4)$$

can be really large if the mean value \bar{x} is small, as is the case when we sum over large numbers of opposite sign.

Generating random numbers: The function

```
numpy.random.normal(mean, sigma, n)
```

returns a sequence of length `n` of values drawn randomly from a normal distribution with mean `mean` and standard deviation `sigma`.

Reading data from a textfile: The command

```
a = numpy.loadtxt('filename.txt')
```

will read data in the file `'filename.txt'` into the array `a`. Note, this assumes the textfile is in the same directory as the python script.

Questions

Note, each of the following 3 questions is equally weighted in the marking scheme.

1. Exploring numerical issues with standard deviation calculations

- (a) Write pseudocode to test the relative error found when you estimate the standard deviation using the two formulas (1) and (2), treating the `numpy` method

```
numpy.std(array, ddof=1)
```

as the correct answer. The input for this calculation will be a supplied dataset consisting of a one-dimensional array of values that is read in using `numpy.loadtxt()`. *Note: The relative error of a value x compared to some true value y is $(x - y)/y$.* In implementing (2), you will need to account for the possibility that this approach could result in taking the square root of a negative number. (Stop and think about why this check is necessary.) You can implement a stopping condition if this is the case, or print a warning.

Submit the pseudo-code.

- (b) Now, using this pseudocode, write a program that uses (1) to calculate the standard deviation of Michelsen's speed of light data (in 10^3 km s^{-1}), which is stored in the file `cdata.txt`, and which was taken from John Baez's (U.C. Riverside) website: http://math.ucr.edu/home/baez/physics/Relativity/SpeedOfLight/measure_c.html

Calculate the relative error with respect to `numpy.std(array, ddof=1)`. Now do the same for eqn. (2). Which relative error is larger in magnitude?

Submit the code, printed output, and written answer. Also include the file `cdata.txt`, separate from your python files and lab report file, in your submission.

- (c) To explore this question further, we will evaluate the standard deviation of a sequence with a predetermined sample variance. Generate two normally distributed sequences, one with

`mean, sigma, n = (0., 1., 2000)`

and another one with

`mean, sigma, n = (1.e7, 1., 2000),`

with the same standard deviation but a larger mean. Then evaluate the relative error of (1) and (2), compared to the `numpy.std` call. How does the relative error behave for the two sequences?

Now that you have investigated a few cases, can you explain the difference in the errors in the two methods, both for these distributions and for the data in Q1b?

Submit printed output and written answer.

- (d) Can you think of a simple workaround for the problems with the one-pass method encountered here? Try this workaround and see if it fixes the problem.

Submit pseudo-code, code, printed output and written answer.

2. Exploring roundoff error

The idea of this exercise is to explore the effects of roundoff error for a polynomial calculated a couple of ways. Consider $p(u) = (1 - u)^8$ in the vicinity of $u = 1$. Algebraically, this is equivalent to the following expansion:

$$q(u) = 1 - 8u + 28u^2 - 56u^3 + 70u^4 - 56u^5 + 28u^6 - 8u^7 + u^8. \quad (5)$$

But numerically, p and q are not exactly the same.

- (a) On the same graph, plot $p(u)$ and $q(u)$ very close to $u = 1$, for example picking 500 points in the range $0.98 < u < 1.02$. Which plot appears noisier? Can you explain why?

Submit your graph

- (b) Now plot $p(u) - q(u)$, and the histogram of this quantity $p(u) - q(u)$ (for u near 1). Do you think there is a connection between the distribution in this histogram and the statistical quantity expressed in equation (3)? To check, first calculate the standard deviation of this distribution (you can use the `std` function in `numpy`). Then calculate the estimate obtained by using equation (3), with $C = 10^{-16}$. State how you calculated the other terms in (3). *Hint: We are looking for order of magnitude consistency here, do not worry about $O(30 - 50\%)$ differences.*

Submit your plots and written answer

- (c) From equation (4) above, show that for values of u somewhat greater than 0.980 but less than 1.0, the error is around 100%. Verify this by plotting or printing out $\text{abs}(p-q)/\text{abs}(p)$ for u starting at $u = 0.980$ and increasing slowly up to about $u = 0.984$ (it might be different on different computers). This fractional error quantity is noisy and diverges quickly as u approaches 1.0, so you might need to plot or print several values to get a good estimate of the values of u at which the error approaches 100%.

Submit your plots and written answer

- (d) Roundoff error doesn't just apply to series, it also comes up in products and quotients. For the same u near 1.0 as in parts a-b, calculate the standard deviation (error) of the numerically calculated quantity $f = u**8/((u**4)*(u**4))$. This quantity will show a range of values around 1.0, with roundoff error. You can get a sense of the error by plotting $f-1$ versus u . Compare this error to the estimate in equation (4.5) on page 131 of the textbook (don't worry about the factor of $\sqrt{2}$).

Submit your plot(s) and written answer

3. Trapezoidal and Simpson's rules for integration

We seek to evaluate the integral

$$\int_0^1 \frac{4}{1+x^2} dx. \quad (6)$$

- (a) What is the exact value of this integral?
- (b) For $N = 4$ slices, compare the value you obtain when using the Trapezoidal vs. Simpson's rule, and with the exact value.
- (c) For each method (Trapezoidal and Simpson), how many slices do you need to approximate the integral with an error of $O(10^{-9})$? We are only looking for a rough estimate for the number of slices for each method (i.e., $N = 2^n$, with $n = 2, 3, 4, \dots$, and N or n being the answer). For this question, do it in a

“dumb” way: increase the number of slices until you hit the mark. How long does it take to compute the integral with an error of $O(10^{-9})$ for each method?

Note: the integration is fast in both methods. In order to get an accurate timing, you may want to repeat the same integration many times over. Recall that we employed a timing method in the first week’s informal lab exercise, but you are free to choose another fancier method.

- (d) Adapt the “practical estimation of errors” of the textbook (Section 5.2.1, page 153) to the trapezoidal method *only* to obtain the error estimation for $N_2 = 32$ (using $N_1 = 16$).
- (e) Why wouldn’t the practical estimation method work with the Simpson’s rule *in our particular case*? How (in a few words; no need to implement it) would we need to adapt the practical estimation of errors method for it to work with the Simpson’s rule *in our particular case*?

For the whole exercise on this integral, submit your pseudo-code, code, printed outputs, and written answers.

Guidelines for evaluation of lab assignments

Collaboration policy

Discussions of questions, ideas on how to approach the problems, and checking if you got the same answers as classmates are the types of collaboration that are allowed. However, **do not copy solutions found from any source** (e.g. other students, websites, textbooks, etc.) This is **plagiarism** and is a **serious academic offence** and we will take appropriate measures. If you do use previous material from a reference, you must cite that reference.

Submitting your work

All labs will be posted and submitted using the *Assignments* tab on Quercus. If you upload your submission early, catch a mistake, and upload a new version, we only mark the latest version.

Each lab will be explicit about what needs to be handed in. You **must** hand in your assignments in two parts using the following naming conventions and formats.

1. python scripts (.py files).

- The grader should be able to run the code on their computer, on a reasonably recent version of python 3 (minimum python 3.6) with the Anaconda distribution, and reproduce your solution.
 - **One main .py file per question.** Combine the scripts for different parts of a question (e.g. parts a and b and c of a Question 2), into a single file. Do not combine questions (e.g. Questions 1 and 2) into a single file.
 - Any input files required for the above python file to run properly must also be included, even if these input files were given out as part of the lab question.
 - If pieces of code were given for you to use as part of the lab question, these pieces must be included in your main .py file (the markers won't add them on to your code).
 - The **name of the main .py file** must be in the following **format**: Lab<lab number>_Q<question number>.py
 - **No jupyter notebooks** (.ipynb files)
- **No interactive components** (e.g. don't require the user to enter input from the keyboard). The .py files must run from the command line.
 - Unless otherwise indicated in the question, the code should not import any libraries besides numpy, scipy, and matplotlib.

2. Everything else, in a single main report file.

- This is the file that the markers will look at first. It must be **.pdf or .doc(x) format**. It contains explanatory notes, pseudocodes (if asked in the lab questions), figures, analysis, etc.
- At the **top** of the file, include your **name**

- Solutions should be typed, unless there are numerous math expressions, in which case you can write those out by hand and submit a readable scan of those equations embedded in your pdf rather than type them out.

Evaluation

Evaluation will consider correctness and quality of solutions. “Good quality” requires being clear enough that another person can quickly and easily understand the problem being solved, the solution approach, and the solution itself.

Main report:

- Solutions should be complete and well documented, with background and solution procedure. They should also be as concise as possible.
- Plots should be clear, well labelled, and well described.
- If your solution is correct, but not clear enough for the marker to understand quickly, you will lose some marks. Unreadable plots or explanations are worthless, in physics as well as in any discipline!

Code:

- Should have a clear logical structure
- Should be perform efficiently (although we don’t expect you to know how to totally optimize performance)
- Include comments. The purpose of comments is so that someone unfamiliar with the specific question (but familiar with programming) can understand how the program works. For example:
- At the beginning of your code, explain its purpose. Also specify if it requires any external function calls, and specify its output (e.g. a single value, a plot of what in particular, an output file with what data...)
- State what variables/constants represent (if not obvious from their names) and units of physical quantities
- Specify mathematical operations for what certain lines are calculating, or what numerical methods are implemented by certain blocks

Lab grading rubric: We will usually use a basic 10-point evaluation rubric for each question, with fractional points permitted. There are two cases, depending on whether your programs are included in the grading. (Please note that even if you are required to hand in your programs, they will not always be graded.) The rubric below is only indicative of a typical lab question.

Case 1: Student is not required to hand in computer program with solution.	Correctness (6/10)	Solution is missing/completely incorrect (0) or complete/perfect (6).
	Presentation (4/10) Solution is very poor quality (0) or very well presented (4).	<ul style="list-style-type: none"> • Problem restated concisely in student’s words. • Solution procedure outlined, including pseudocode if required. (Pseudocode can contribute to a good solution, even if it is not required.) • Text output is clearly presented (column format), plots are high quality (title, labels, captions, legends, clarity of graphics, etc).

Case 2: Student is required to hand in computer program with solution.	Correctness (6/10)	Solution is missing/completely incorrect (0) or complete/perfect (6).
	Presentation (2/10) Solution is very poor quality (0) or very well presented (2).	<ul style="list-style-type: none"> · Problem restated concisely in student's words. · Solution procedure outlined, including pseudocode if required. (Pseudocode can contribute to a good solution, even if it is not required.) · Text output clearly presented (column format), plots are high quality (title, labels, captions, legends, clarity of graphics, etc).
	Code (2/10) Code provided is very poor quality (0) or excellent quality (2)	<ul style="list-style-type: none"> · Program is well structured, logic is easy for someone to understand. · Program is well documented with comments and easy to understand variable names. · Test of program is successful (occasionally used at instructor's discretion).